

# LEARNING VELOCITY PRIOR-GUIDED HAMILTONIAN-JACOBI FLOWS WITH UNBALANCED OPTIMAL TRANSPORT

Anonymous authors

Paper under double-blind review

## ABSTRACT

The connection between optimal transport (OT) and control theory is well established, most prominently in the Benamou–Brenier dynamic formulation. With quadratic cost, the OT problem can be reframed as a stochastic control problem in which a density  $\rho_t$  evolves under a controlled velocity field  $v_t$  subject to the continuity equation  $\partial_t \rho_t + \nabla \cdot (\rho_t v_t) = 0$ . In this work, we introduce a velocity prior into the continuity equation and derive a new Hamilton–Jacobi–Bellman (HJB) formulation to learn dynamical probability flows. We further extend the approach to the unbalanced setting by adding a growth term, capturing mass variation processes common in scientific domains such as cell proliferation and differentiation. Importantly, our method requires training only a single neural network to model  $v_t$ , without the need for a separate model for the growth term  $g_t$ . Finally, by decomposing the velocity field as  $v_{\text{total}} = v_{\text{prior}} + v_{\text{corr}}$ , our approach is able to capture complex transport patterns that prior methods struggle to learn due to the curl-free limitation.

1 INTRODUCTION

From flow matching (FM) to action matching (AM), learning transport maps between distributions has been widely explored in recent years (Lipman et al., 2022; Albergo and Vanden-Eijnden, 2022; Liu et al., 2022; Neklyudov et al., 2023a). *Flow Matching (FM)* (Lipman et al., 2022) learns a time-dependent velocity field  $u_t$  that pushes  $\rho_0$  to  $\rho_1$  and can realize highly expressive transport paths; however, the original FM with independent coupling between source and target does not guarantee *least action* by minimizing the kinetic energy in the Benamou–Brenier sense. Instead, it trains  $u_t$  to match conditional expectations of displacement vectors under a chosen interpolation scheme, which may yield non-optimal flows.

*Action Matching* (AM) (Neklyudov et al., 2023a) addresses this by parameterizing a scalar potential  $s_t$  whose gradient  $\nabla s_t$  induces the transport, aligning with the optimality conditions of OT and yielding lower kinetic energy than unconstrained FM. The price is reduced expressiveness:  $\nabla s_t$  is *curl-free*, so AM cannot directly represent rotational or cyclic dynamics that are common in scientific domains. From the Helmholtz decomposition perspective (Neklyudov et al., 2023a), any vector field  $u_t^*$  can be written as  $u_t^* = \nabla s_t^* + w_t$  with  $w_t$  divergence-free (Ambrosio et al., 2005, §8.4.2). Under this lens, AM retains only the gradient component and discards  $w_t$ , explaining both its energy efficiency and its inability to encode rotations and cycles.

In this paper, we seek a middle ground – expressive like FM, energy-aware like AM – by introducing a velocity prior  $v_{\text{prior}}$  and learning only the residual potential. We note that even compared with energy-aware FM variants such as OT-CFM(Pooladian et al., 2023; Tong et al., 2023a), our approach achieves better energy efficiency, as demonstrated in upper Table 2. Specifically, we decompose the velocity field as  $v_{\text{total}}(t, x) = v_{\text{prior}}(t, x) + \nabla s_t(x)$ . Here  $v_{\text{prior}}$  captures known rotational dynamics or domain-specific effects such as RNA velocity in single-cell biology, while  $\nabla s_t$  accounts for the OT-consistent gradient component. We train  $s_t$  by minimizing a modified Hamilton–Jacobi residual that incorporates the prior, together with boundary terms that ensure  $\rho_0 \rightarrow \rho_1$ . This *residualized* design preserves OT optimality conditions for the learned component, improves interpretability,

and injects inductive bias without paying the kinetic-energy cost of unconstrained original FM. We name our approach the Velocity Prior Hamiltonian-Jacobi Flow (VP-HJF).

**Motivation in practice** In domains like single-cell biology and physical systems with known drifts, accurate priors are available yet incomplete and mass change such as cell proliferation and decay is ubiquitous. VP-HJF exploits these priors to encode hard-to-learn structure. The residual potential learns both the correction that the prior can not explain and the mass changes through the growth term. This yields a compact and interpretable alternative to fully free vector-field models, particularly effective when local supervision is noisy but prior knowledge is rich.

## 2 BACKGROUND

**Dynamical Optimal Transport** Beyond the classic static Monge–Kantorovich formulation in OT (Ambrosio et al., 2005; Villani et al., 2008), there exists a dynamical formulation known as the Benamou–Brenier problem which links OT with PDEs by representing the  $W_2$  distance as the minimum kinetic energy where  $\rho_t$  is density and  $v_t$  is a velocity field with boundary conditions:  $\rho|_{t=0} = \rho_0$ ,  $\rho|_{t=1} = \rho_1$ , (Benamou and Brenier, 2000):

$$W_2^2(\rho_0, \rho_1) = \inf_{\rho_t, v_t} \int_0^1 \int \frac{1}{2} \|v_t(x)\|^2 \rho_t(x) dx dt, \quad \partial_t \rho_t + \nabla \cdot (\rho_t v_t) = 0. \quad (1)$$

**Unbalanced Optimal Transport** When total mass change over time such as following a growth-decay process in biology, we add a growth rate  $g_t(x)$  term to the continuity equation to incorporate the weight changes (Chizat et al., 2018):

$$\partial_t \rho_t(x) + \nabla \cdot (\rho_t(x) v_t(x)) = g_t(x) \rho_t(x), \quad \rho|_{t=0} = \rho_0, \quad \rho|_{t=1} = \rho_1, \quad (2)$$

The Wasserstein–Fisher–Rao distance with scale  $\delta > 0$  is defined as the minimal action balancing transport cost and mass change:

$$\text{WFR}_\delta^2(\rho_0, \rho_1) = \inf_{\rho, v, g} \int_0^1 \int \left( \frac{1}{2} \|v_t(x)\|^2 + \frac{\delta^2}{2} g_t(x)^2 \right) \rho_t(x) dx dt, \quad \text{s.t. Eq.2.} \quad (3)$$

**Hamilton–Jacobi–Bellman (HJB)** We recall the classical connection between optimal control and Hamilton–Jacobi (HJ) theory. Consider a deterministic control system with state  $x(t) \in \mathbb{R}^d$ , control  $u(t)$ , dynamics  $\dot{x} = f(x, u, t)$ , running cost  $L(x, u, t)$ , and terminal cost  $\psi(x)$ . The *value function*

$$V(t, x) = \inf_{u(\cdot)} \left\{ \int_t^1 L(x(s), u(s), s) ds + \psi(x(1)) \right\}$$

gives the minimal cost-to-go from  $(t, x)$  under admissible controls. It is well known that  $V$  solves the Hamilton–Jacobi–Bellman (HJB) PDE

$$\partial_t V(t, x) + H(x, \nabla V(t, x), t) = 0, \quad V(1, x) = \psi(x),$$

where the Hamiltonian is

$$H(x, p, t) := \inf_u \left\{ L(x, u, t) + p^\top f(x, u, t) \right\}, \quad p = \nabla V(t, x).$$

**Action Matching (AM)** AM fits a scalar potential  $s_\theta$  to learn a energy-minimizing flow between distributions by minimizing the (un)balanced HJB residuals.

$$\mathcal{L}_{\text{uAM}} = \int_0^1 \mathbb{E}_{x \sim \rho_t} \left[ \partial_t s_\theta(t, x) + \frac{1}{2} \|\nabla_x s_\theta(t, x)\|^2 + \frac{1}{2} s_\theta^2(t, x) \right] dt, \quad (4)$$

with boundary constraints as:  $\mathbb{E}_{x \sim \rho_0}[s_0(x)] - \mathbb{E}_{x \sim \rho_1}[s_1(x)]$ .

108    **3 METHODOLOGY**  
 109

110    We introduce a velocity-prior guided approach, the *Velocity Prior Hamiltonian–Jacobi Flow (VP-HJF)*, to solve the unbalanced optimal transport problem under the Wasserstein–Fisher–Rao (WFR) metric (Eq. 2, 3). In contrast to prior approaches that fit two separate networks—one for transport and one for growth (Zhang et al., 2024; Wang et al., 2025), our method trains a single neural network. Following (Neklyudov et al., 2023a), we can represent both the transport velocity field and the growth term through a single scalar potential.

116    **Proposition 3.1** (Neklyudov et al., 2023a, Prop. 3.3). *Suppose we have a continuous dynamic flow  
 117    with density  $\rho_t$ . Under mild conditions, there exists a unique scalar potential function  $\hat{s}_t(x)$  such  
 118    that the unbalanced continuity equation (2) is satisfied, with the velocity field and growth function  
 119    given by  $v_t^*(x) = \nabla \hat{s}_t(x)$ ,  $g_t^*(x) = \hat{s}_t(x)$ .*

120    Building on Proposition 3.1, we reduce the WFR problem to learning a single model and incorporate  
 121    problem-specific dynamics through a simple velocity decomposition. Specifically, we decompose  
 122    the velocity field into two parts: a known velocity prior and a learnable corrective velocity field  
 123    component:

$$v_{\text{total}}(t, x) = v_{\text{prior}}(t, x) + v_{\text{corr}}(t, x), \quad (5)$$

124    where  $v_{\text{prior}}$  encodes domain knowledge (e.g. translations, rotations, RNA velocity), and  $v_{\text{corr}}$  is  
 125    the data-driven corrective component. In this way, the prior captures coarse dynamics while  
 126    the model focuses on refinements such as correcting the residual transport and learning mass imbal-  
 127    ance that the prior cannot explain. In essence, our approach improves interpretability and reduces  
 128    the learning complexity through adding prior knowledge of the velocity field  $v_{\text{prior}}$  – leaving the  
 129    learnable velocity field  $v_{\text{corr}}$  simpler learning tasks compared with other generative modeling meth-  
 130    ods of learning the entire velocity field  $v_{\text{total}}$ . Intuitively, our approach pays kinetic cost only for  
 131    the *correction* to the prior drift and for the mass growth-decay component, making learning more  
 132    efficient.

134    We can now define our velocity-prior guided unbalanced OT problem under the least-action principle  
 135    as:

136    **Definition 3.2.** Consider the following least-action objective with  $\delta=1$  and subject to the unbalanced  
 137    *velocity-prior guided* continuity equation :

$$\mathcal{A}(\rho, v, g) = \int_0^1 \int \left( \frac{1}{2} \|v(t, x)\|^2 + \frac{1}{2} g(t, x)^2 \right) \rho_t(x) dx dt, \quad (6)$$

$$\text{s.t. } \partial_t \rho_t = -\nabla \cdot (\rho_t (v_{\text{prior}} + v_{\text{corr}})) + g_t \rho_t, \quad \rho|_{t=0} = \rho_0, \quad \rho|_{t=1} = \rho_1. \quad (7)$$

143    Note that in our method we do *not* optimize over  $\rho$  directly. Instead,  $\rho_t$  is *induced* by a parametric  
 144    flow  $\Phi_t^\theta$  via  $\dot{x} = v_{\text{prior}}(t, x) + \nabla s_\theta(t, x)$  and defined as  $\rho_t^\theta = (\Phi_t^\theta)_\# \rho_0$ .

146    **Prior-guided HJB residual** Since solving for the minimum-action problem in primal form in  
 147    Definition 3.2 is intractable, we turn to its dual formulation. The key derivation step is to introduce  
 148    a scalar potential  $s(t, x)$  as the Lagrange multiplier for the prior-guided continuity equation and  
 149    applying the Fenchel–Young inequalities to the velocity field and growth term. We then obtain the  
 150    following dual lower bound (see Appendix A for details):

$$\begin{aligned} \mathcal{A}(\rho, v, g) &\geq \mathbb{E}_{\rho_0(x)}[s_0(x)] - \mathbb{E}_{\rho_1(x)}[s_1(x)] \\ &\quad - \int_0^1 \int \rho_t(x) \left( \partial_t s + \frac{1}{2} \|\nabla s\|^2 + \nabla s \cdot v_{\text{prior}} + \frac{1}{2} s^2 \right) dx dt. \end{aligned} \quad (8)$$

155    The bound is tight point-wise if and only if when we choose the primal variables as

$$v_{\text{corr}}(t, x) = \nabla_x s(t, x), \quad g(t, x) = s(t, x),$$

158    which shows that  $s$  can simultaneously control both the *corrective* transport  $\nabla_x s$  and the local  
 159    growth  $s$ . We can then plug these back to the continuity equation to get the optimal particle dynamics  
 160    and their log-weights evolve as

$$\dot{x}(t) = v_{\text{total}}(t, x) = v_{\text{prior}}(t, x) + \nabla_x s(t, x), \quad \frac{d}{dt} \log w(t) = s(t, x(t)).$$

162 **Corollary 3.3** (HJB residual objective). *Motivated by the duality form, we propose to parameterize*  
 163  *$s_\theta(t, x)$  with a neural network and define the velocity-prior guided HJB residual as:*

$$165 \quad r_\theta(t, x) := \partial_t s_\theta(t, x) + \frac{1}{2} \|\nabla_x s_\theta(t, x)\|^2 + \nabla_x s_\theta(t, x) \cdot v_{\text{prior}}(t, x) + \frac{1}{2} s_\theta(t, x)^2. \quad (9)$$

166 Then minimizing

$$168 \quad \mathcal{L}_{\text{HJB}}(\theta) = \mathbb{E}_{\rho_0(x)}[s_0] - \mathbb{E}_{\rho_1(x)}[s_1] + \int_0^1 \mathbb{E}_{\rho_t(x)}[w(t, x)r_\theta(t, x)^2] dt \quad (10)$$

171 drives  $s_\theta$  toward dual feasibility. Note that in practice, we use squared residual to prevent positive  
 172 and negative values from cancellation and add a importance weight  $w(t)$  trick to reduce variance.

174 **Importance Reweighting** The squared HJB residual can be dominated by a few high-variance  
 175 outliers (rare cells, sharp local flows), which destabilizes training. To ensure training stability and  
 176 preventing these extreme high residual outliers, we adopt a simple batch-wise importance reweighting  
 177 that down-weights large residuals. For a mini-batch  $\{(t_i, x_i)\}_{i=1}^B$ , let  $r_i = |r_\theta(t_i, x_{t,i})| + \varepsilon$   
 178 and with a temperature  $\tau > 0$ . Then for each sample, the weight is inversely proportions to a  
 179 temperature-shaped residual as  $\tilde{w}_i \propto r_i^{-\tau}$ . Thus, larger residuals get smaller weight, which re-  
 180 duces variance while keeping the update focused and stable.

181 **Theorem 3.4** (Prior-guided HJB optimality). *Suppose that the HJB residual defined in corollary.3.3*  
 182 *satisfies  $r_\theta(t, x) = 0$  for  $\rho_t$ -a.e on  $[0, 1] \times \mathbb{R}^d$ , and the boundary constraints hold, then  $(\rho_t, v_{\text{total}}, g_\theta)$*   
 183 *satisfies the unbalanced continuity equation and the WFR optimality conditions in Definition 3.2. In*  
 184 *particular, the learned corrective field  $v_{\text{corr}}^* = \nabla_x s_\theta$  and growth  $g^* = g_\theta$  satisfy the optimality*  
 185 *conditions.*

186 While the HJB residual enforces local optimality conditions, it does not guarantee that the terminal  
 187 distribution  $\rho_1^\theta$  matches with the target  $\rho_1$ . To bridge this gap, we design a two-part reconstruction  
 188 objective: (i) a density matching term through the sliced Wasserstein between the predicted  $\hat{\rho}_1$  and  
 189  $\rho_1$ , and (ii) a mass term aligning the global log-mass ratio. These two terms directly calibrates the  
 190 terminal distribution's *shape* and *mass*, complementing the HJB residual.

191 **Reconstruction loss** To align the terminal distribution in *shape*, we use a sliced Wasserstein ob-  
 192 jective. Let  $\hat{\rho}_1$  be the empirical terminal distribution learned from our model, and  $\rho_1$  be the ground  
 193 truth target distribution.  $\theta_\ell$  is a random projection sampled from  $\theta_\ell \sim \text{Unif}(\mathbb{S}^{d-1})$  for  $\ell = 1, \dots, L$ ,  
 194  $\hat{X}_1$  are predicted samples, and  $Y$  are ground-truth samples, the sliced Wasserstein loss is defined as:

$$197 \quad \text{SW}_2^2(\hat{\rho}_1, \rho_1) \approx \frac{1}{L} \sum_{\ell=1}^L W_2^2(\langle \theta_\ell, \hat{X}_1 \rangle, \langle \theta_\ell, Y \rangle), \quad (11)$$

199 where  $W_2^2$  on  $\mathbb{R}$  is the 1D Wasserstein distance computed by sorting projections.

201 To capture *global mass change*, we integrate particle log-weights along the learned dynamics. With  
 202 WFR scale  $\delta = 1$  and a mini-batch of size  $B$ . At  $t = 0$ , we initialize the weights as  $\log w_i(0) = 0$ ,  
 203 and follow the below characteristic ODEs to calculate the predicted terminal mass ratio at  $t = 1$   
 204 relative to  $t = 0$ :

$$206 \quad \frac{d}{dt} \log w_i(t) = s_\theta(t, x_i(t)), \quad \hat{r} = \frac{1}{B} \sum_{i=1}^B e^{\log w_i(1)} = \{\log w_i(1)\}_{i=1}^B \quad (12)$$

208 We penalize against the ground truth ratio  $r$  via its logarithmic form to ensure stability and enforces  
 209 the WFR consistency:  $\mathcal{L}_{\text{mass.}} = (\log \hat{r} - \log r)^2$ . Our final reconstruction loss combines the two  
 210 components with tunable coefficients:

$$213 \quad \mathcal{L}_{\text{recon}} = \lambda_{\text{sw}} \text{SW}_2^2(\hat{\rho}_1, \rho_1) + \lambda_{\text{mass}} (\log \hat{r} - \log r)^2 \quad (13)$$

215 with  $\lambda_{\text{sw}} > 0$  and  $\lambda_{\text{mass}} > 0$ . In practice we use  $L \in [64, 512]$  random projections, and choose  
 $\lambda_{\text{mass}} \in [0.1, 1]$  to calibrate mass without overpowering other terms.

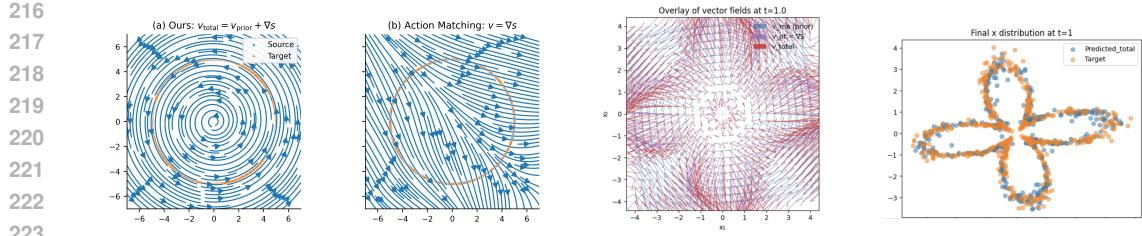


Figure 1: **Left** (first two): Ours correctly learned the rotating dynamic while AM failed. **Middle**: Vector field (red) bending away from  $v_{\text{prior}}$  (blue) to form petal shapes. **Right**: Predicted target distribution matches with ground truth

**Total objective** Putting the pieces together, our total training loss is

$$\min_{\theta} \mathcal{L}(\theta) = \underbrace{\mathbb{E}_{x \sim \rho_0}[s_0(x)] - \mathbb{E}_{x \sim \rho_1}[s_1(x)] + \int_0^1 \mathbb{E}_{x \sim \rho_t}[w(t, x) r_\theta(t, x)^2] dt}_{\mathcal{L}_{\text{HJB}}} + \underbrace{\lambda_{\text{sw}} \text{SW}_2^2(\hat{\rho}_1, \rho_1) + \mathcal{L}_{\text{mass}}}_{\mathcal{L}_{\text{recon}}} \quad (14)$$

where  $r_\theta(t, x)$  is the HJB residual,  $w(t, x)$  are the nonnegative weights, and  $\hat{\rho}_1(\theta)$  is the terminal distribution obtained by pushing  $\rho_0$  through the learned dynamics.

## 4 RELATED WORKS

**Physics-constrained approach** Existing works (Koshizuka and Sato, 2022; Neklyudov et al., 2023b; Tong et al., 2020) add a *potential-based prior*  $V_t(x)$  to the HJ equation to incorporate prior knowledge for trajectory inference. In (Neklyudov et al., 2023b), this yields  $\partial_t s + \frac{1}{2} \|\nabla s\|^2 + V_t + \frac{1}{2\delta^2} s^2 = 0$  and the conservative second-order law  $\ddot{X}_t = -\nabla V_t(X_t)$ . Meanwhile, our method is the *velocity-based prior* approach, the known velocity field  $v_{\text{prior}}$  enters as a drift inside the continuity equation, which produces the cross-term  $\nabla s \cdot v_{\text{prior}}$ . Since potential-based priors are curl-free, it can not represent rotational flows as in (Neklyudov et al., 2023a); By contrast, our data-driven formulation uses a flexible measured vector fields directly as  $v_{\text{prior}}$ , which can be considered as a "free drift" and learns only minimal optimal corrections and growth.

**Other trajectory inference approaches.** Trajectory inference has also advanced through *flow matching* and *Schrödinger bridge* methods, which scale effectively to high-dimensional data. Recent SB variants further improve performance on single-cell datasets include (Huguet et al., 2022; Tong et al., 2023b). For unbalanced settings, variational and regularized UOT methods such as DeepRUOT, VRUOT, and VGFM directly learn transport dynamics and growth from snapshot data (Zhang et al., 2024; Sun et al., 2025; Wang et al., 2025). In particular, VRUOT also uses a single network to model both the velocity field and the growth term, but it did not incorporate prior known knowledge like ours.

## 5 EXPERIMENTS

### 5.1 SYNTHETIC DATASET

**Balanced case - Rotating Ring** First, we show a case when utilizing the velocity prior is crucial in learning the correct velocity field where curl-free methods like AM fails. We tested on a 2D rotating ring dataset where the points on the ring (source) are rotated by a fixed angle  $\theta$  (target). The velocity prior is defined as  $v_{\text{prior}} = \omega J x$ , where  $J$  is the skew-symmetric rotation matrix and  $x \in \mathbb{R}^2$ , so this becomes  $Jx = [0 \ 1; -1 \ 0][x_1 \ x_2]^\top = (-x_2, x_1)^\top$ . The task for our model  $v_{\text{ot}}$  is to

| Method                    | $W_2$        | Control action<br>( $\nabla s$ ) | Total kinetic<br>( $v_{\text{total}}$ ) |
|---------------------------|--------------|----------------------------------|---|
| Flow Matching             | 0.582        | 18.369                           | 18.369                                  |
| OT-FM                     | <b>0.402</b> | 18.707                           | 18.707                                  |
| VP-HJF (Ours)             | 0.577        | <b>0.624</b>                     | <b>17.955</b>                           |
| Prior-only ( $\alpha=1$ ) | 1.351        | 0                                | 36.250                                  |

| Method                                    | $r_{\text{pred}}$ | $ \Delta \log r $ | Rel. % err. |
|---|-------------------|-------------------|-------------|
| VP-HJF (ours)                             | 1.356             | 0.044             | 4.30%       |
| VP-HJF (w/o $\mathcal{L}_{\text{mass}}$ ) | 0.899             | 0.455             | 36.56%      |
| Unbalanced AM                             | 0.938             | 0.413             | 33.80%      |

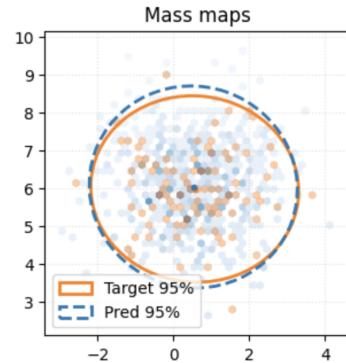


Figure 2: **Left:** Upper table on Least Action Comparison for Balanced Gaussian Translation; **Bottom table:** Effect of the mass loss  $\mathcal{L}_{\text{mass}}$ .  $r_{\text{pred}}$  is the predicted mass ratio , the absolute log error  $|\Delta \log r| = |\log r_{\text{pred}} - \log r_{\text{true}}|$ , and the relative error (%), with  $r_{\text{true}} = 1.418$ . **Right:** Ellipses shows the mean/covariance of our approach matches with Target for Unbalanced Gaussian Translation; Inside, each bin indicates weights. see Appendix B

learn the residual correction after given the prior rotation knowledge, such as ensuring the boundary condition by aligning the mismatched source and target density or correcting the radial drift by push the points inwards or outwards,etc. In Figure 1 left (b), we show that since AM has the curl-free limitation, without a prior, its model  $\nabla s_t(x)$  failed to represent pure rotation where the streamlines cut through the circle. In Figure 1 left (a), the streamlines from our method form a circular flow indicating that  $v_{\text{prior}}$  gives the model the correct inductive bias.

**Diverging Petal** We created a curved and rotated petal-shape dataset to test our method on diverging multi-trajectory paths. The source is a gaussian distribution concentrated in the center and  $v_{\text{prior}=\omega J_x}$  the rotation dynamic defined as before. Our task is to learn  $v_{\text{petal}}$ , a radial and angle-dependent term that push outward the points along the radius with different speed depend on the angle  $\theta = \text{atan2}(y, x)$ . We have

$$v_{\text{petal}}(x) = s(\theta) \hat{r}, \quad \hat{r} = \frac{x}{\|x\|}, \quad s(\theta) = \max(0, b + a \cos(k\theta)) \quad (15)$$

Compared with the petal shape appeared in AM and MioFlow (Huguet et al., 2022), the underlying dynamic flow in our example is harder to learn, where the former one has a straight-axis aligned radial expansion as  $r(x) = |x_1| + |x_2|$  with the gradient of  $r(x)$  being a piece-wise constant and curl-free. Figure 1 middle shows that our petal shape matches with the target shape. Figure 1 right shows that the vector field (red arrows) are bending away from pure rotation (blue arrows) to align strongly with the petal shape by pushing the mass outward.

**Balanced case - Gaussian Translation** In this experiment, we compare the least-action or energy costs across different methods. As shown in Table 2(upper), our proposed VP-HJF achieves Wasserstein distance  $W_2$  accuracy on par with FM Lipman et al. (2022) while OT-FM Tong et al. (2023a); Pooladian et al. (2023) attains the lowest Wasserstein distance. By contrast, both FM and OT-FM must learn the entire velocity field  $u_t = \nabla s + w_t$ , resulting in a much larger control action. This demonstrates that VP-HJF leverages the structured prior effectively, where the prior dynamics carry most of the transport, and the learned correction  $\nabla s$  makes adjustments. To verify that our improvement is not solely due to a strong velocity prior  $v_{\text{prior}}$  itself, we also report a prior-only baseline. The prior alone shows moderate accuracy but with high kinetic cost, whereas our method is able to balance both accuracy and energy efficiency. Additionally, Fig. 1 highlights that incorporating the velocity prior is crucial in the rotational dynamics setting, as well as in a multi-trajectory example where paths diverge into a petal shape.

**Lotka–Volterra with growth.** We model the prey and predator densities  $x_1(t), x_2(t)$  by a first-order nonlinear ODE,

$$\dot{x}_1(t) = \alpha x_1(t) - \beta x_1(t) x_2(t), \quad \dot{x}_2(t) = -\gamma x_2(t) + \delta x_1(t) x_2(t),$$

| $\Delta \log r$ (abs) ↓ | Clean | Gaussian noise $\alpha$ | Scale $c$ | Wrong coeff |
|-------------------------|-------|-------------------------|-----------|-------------|
|                         | 0.5   | 0.75                    | 0.50 0.75 | medium hard |
| <b>VP-HJF (ours)</b>    | 0.04  | 0.17                    | 0.02 0.11 | 0.21 0.37   |
| Prior-only <sup>†</sup> | 0.35  | 0.35                    | 0.35 0.35 | 0.35 0.35   |

Table 1: Ablations on prior robustness (LV). The absolute  $\Delta \log r := |\log \hat{r} - \log r|$ . Prior-only corresponds to the ground truth mass ratio (mass-conserving). Wrong coeff medium uses  $\alpha = 1.3, \beta = 0.8, \gamma = 1.2, \delta = 0.7$  and hard uses  $\alpha = 1.5, \beta = 0.6, \gamma = 1.4, \delta = 0.6$

---

**Algorithm 1** Training VP–HJF (Velocity-Prior Hamiltonian–Jacobi Flow)

---

```

335 Require: per time-interval  $t_k$  snapshots  $\{x, t_k, v_k^{prior}\}_{k=0}^K$ , network  $s_\theta(t, x)$ , coefficients
336  $\alpha, \beta, \lambda_{SW}, \lambda_{mass}$ ,
337 1: while not converged do
338 2:   Sample interval index  $k \in \{0, \dots, K - 1\}$ ,
339 3:   Sample adjacent pairs  $(x_k^{(b)}, v_k^{(b)}, x_{k+1}^{(b)}, v_{k+1}^{(b)})_{b=1}^B \sim \{x, t, v_{prior}\}_{k=0}^K$ 
340 4:   Sample normalized  $t \sim (0, 1)$ 
341 5:   HJB residual loss:
342 6:   Compute  $s_k \leftarrow s_\theta(t_k, x_k), s_0 \leftarrow s_\theta(t_0, x_k), s_1 \leftarrow s_\theta(t_1, x_{k+1})$ ;
343 7:   Compute importance weights  $w_b \propto r_{k,b}^{-\tau}$ 
344 8:    $r_k \leftarrow \partial_t s_k + \frac{1}{2} \|\nabla_x s_k\|^2 + \nabla_x s_k \cdot v_{prior}(t_k, x_k) + \frac{1}{2} s_k^2$ 
345 9:    $\mathcal{L}_{HJB} \leftarrow \mathbb{E}_{\rho_{t_0}}[s_0] - \mathbb{E}_{\rho_{t_1}}[s_1] + \frac{1}{B} \sum_{b=1}^B w_b r_{k,b}^2$ 
346 10:  Reconstruction loss:
347 11:  Initialize particles  $x^{(b)} \leftarrow x_k^{(b)}$  and  $\log w^{(b)} \leftarrow 0$  for  $b = 1..B$ 
348 12:  ode rhs =  $[\dot{x} = \nabla_x s_\theta(t, x) + v_{prior}(t, x), \log w = s_\theta(t, x)]$ 
349 13:   $(x_{k+1}^{\text{pred}}, \log w_{k+1}) \leftarrow \text{odeint}(\text{ode rhs}, (x_k, 0), t \in [t_k, t_{k+1}])$ 
350 14:  Sliced Wasserstein loss:  $\mathcal{L}_{SW} \leftarrow \frac{1}{L} \sum_{\ell=1}^L W_2^2(x_{k+1}^{\text{pred}}, x_{k+1})$ 
351 15:  Mass loss:  $\mathcal{L}_{mass} \leftarrow (\log \hat{w}(1) - \log w(1))^2$ 
352 16:   $\mathcal{L}_{recon} \leftarrow \lambda_{SW} \mathcal{L}_{SW} + \lambda_{mass} \mathcal{L}_{mass}$ 
353 17:  Total loss:  $\mathcal{L}(\theta) \leftarrow \alpha \mathcal{L}_{HJB} + \beta \mathcal{L}_{recon}$ 
354 18:  Update  $\theta \leftarrow \theta - \eta \nabla_\theta \mathcal{L}(\theta)$ 
355 19: end while

```

---

where  $\alpha$  is the prey’s intrinsic growth rate,  $\beta$  is the predation rate,  $\gamma$  is the predator’s mortality rate, and  $\delta$  is the predator’s growth rate from consuming prey. To model the population expansion and decay dynamics, we use a simple scalar growth field and evolve local mass via the weight dynamics

$$g(x(t)) = \kappa (x_1(t) - x_2(t)), \quad \frac{d}{dt} \log w(t) = g(x(t))$$

The total mass  $M(t) = \mathbb{E}[w(t)]$  and the ground-truth mass ratio is calculated as  $r_{\text{true}} = M(t)/M(0)$ .

Table 2 (bottom) shows that adding the explicit mass term  $\mathcal{L}_{mass}$  enables our method to closely match  $r_{\text{true}}$  with 4% of relative error while both unbalanced AM and our method without the  $\mathcal{L}_{mass}$  term suffer from a much higher relative error of over 30%. Although AM aligns transport but leaves the *scale* of the scalar potential  $s_\theta$  unconstrained, so the integrated growth  $\int_0^t g(x(t)) dt$  is miscalibrated. By contrast,  $\mathcal{L}_{mass}$  provides endpoint constraint on the mass – yielding better mass dynamics and more aligned with the ground truth mass ratio.

**Ablations on prior robustness** In this section, we conduct various ablation studies to test the robustness of our model. In particular, we modify the quality of  $v_{prior}$  by misspecifying it to different scaling  $v_{prior} = cv$ , adding Gaussian noises  $v_{prior} = v + \eta$  and using the wrong coefficients with various level of difficulties in the LV equation. We compare it with the prior-only approach which

378 Table 2: Comparison on the EB dataset using SWD at the held-out marginals ( $t_1, t_3$ ) and the mean  
 379 over the remaining snapshots. Baseline results are taken from (Theodoropoulos et al., 2025)  
 380

| Method               | SWD $t_1$                         | SWD $t_3$                         | Rest SWD                          |
|----------------------|-----------------------------------|-----------------------------------|-----------------------------------|
| SBIRR                | $0.80 \pm 0.06$                   | $0.91 \pm 0.05$                   | $0.66 \pm 0.07$                   |
| MMFM                 | $0.59 \pm 0.04$                   | $0.76 \pm 0.04$                   | $0.52 \pm 0.07$                   |
| DMSB                 | $0.58 \pm 0.06$                   | $0.54 \pm 0.06$                   | $0.45 \pm 0.04$                   |
| 3MSBM                | $0.48 \pm 0.04$                   | <b><math>0.38 \pm 0.03</math></b> | <b><math>0.36 \pm 0.05</math></b> |
| <b>VP-HJF (ours)</b> | <b><math>0.37 \pm 0.01</math></b> | $0.46 \pm 0.01$                   | $0.50 \pm 0.08$                   |

381 essentially corresponds to the ground-truth mass ratio  $r_{\text{true}}$  as mass is conserved in this case. Table  
 382 1 shows that under global scaling our approach is particularly stable – indicating that the learned  
 383 corrective field  $v_{\text{corr}}$  and growth term calibrate mass well when the prior’s direction is correct up  
 384 to scale. With additive Gaussian perturbations, errors remain modest showing tolerance to stochastic  
 385 mis-specification. The primary failure is the wrong coefficient (hard)case, which reflects the  
 386 sensitivity to severe directional mismatch that cannot be fully compensated without incurring large  
 387 kinetic corrections.

## 388 5.2 REAL-WORLD DATASET

389 In this section, we evaluate our method on two single-cell RNA-seq datasets. Both provide RNA ve-  
 390 locity, which we use as the velocity prior  $v_{\text{prior}}$ . Such priors are common in biological and scientific  
 391 applications beyond single-cell data. Incorporating them introduces an inductive bias that reduces  
 392 learning complexity — our model needs only to learn a corrective flow and growth rather than the  
 393 full dynamics from scratch.

404 **EB scRNA-Seq data** We evaluate cell-trajectory inference on the Embryoid Body (EB) dataset  
 405 of Moon et al. (2019), using the preprocessed release from Koshizuka and Sato (2022); Tong et al.  
 406 (2020). The dataset comprises five snapshots over 27 days, grouped as  $t_0 \in [0, 3]$ ,  $t_1 \in [6, 9]$ ,  
 407  $t_2 \in [12, 15]$ ,  $t_3 \in [18, 21]$ ,  $t_4 \in [24, 27]$ . Leveraging RNA velocity as a prior  $v_{\text{prior}}$  at each snapshot,  
 408 we train a local and shorter trajectory by adopting the *local per-interval* training: at each step we  
 409 sample an adjacent pair  $(t_k, t_{k+1})$  and learn only the transport and growth to move  $\rho_{t_k} \rightarrow \rho_{t_{k+1}}$ .  
 410 This yields more stable gradients and low target variance than enforcing all time points jointly;. For  
 411 details of the training algorithm, see Algorithm 1.

412 We test on 100-dim PCA components feature space and compare with recent works using the multi-  
 413 marginal approach from 3MSBM(Theodoropoulos et al., 2025), SBIRR (Shen et al., 2024), MMFM  
 414 (Rohbeck et al., 2025) and DMSB (Chen et al., 2023). We follow the experiment setup from 3MSBM  
 415 by having  $t = 1, 3$  as the held out sets. Table 2 shows that our method outperforms others at  $t = 1$   
 416 and remains competitive at other time snapshots in sliced Wasserstein distance. Notably, our method  
 417 outperforms SBIRR and MMFM which solve a piecewise Schrödinger Bridges and OT coupling  
 418 instead of a single global optimization with a joint coupling like 3MSBM and DMSB, indicating the  
 419 benefits of using RNA velocity as a local prior with per-interval supervision.

420 **Bone marrow scRNA-Seq data** We evaluate our approach on a real scRNA-seq bone-marrow  
 421 atlas with multiple hematopoietic fates from scVelo (Bergen et al., 2020). Figure 3 (left) shows  
 422 trajectories that emanates from early progenitor regions at  $t=0$  (dark blue) and spread out to other  
 423 branches by closely following the UMAP reference (gray). Figure 3 (right) shows the learned growth  
 424 field  $g_\theta(t, x) = s_\theta(t, x)$  that governs local mass dynamics. This maps shows that our model suc-  
 425 cessfully captured low growth rate in early progenitors and increases as cells enter the active cycling  
 426 and amplification stage. However, we also observe that the high growth rate near some terminal  
 427 regions. While this indicates the model can assign higher growth to specific cell types, a strong  
 428 terminal-phase growth is biologically implausible. For the bone marrow case, mature or exiting cells  
 429 should have near-zero or negative growth. This likely reflect the objective imbalance where trans-  
 430 port terms dominating mass calibration, suggesting mild regularization such as time-smoothness on  
 431  $s_\theta$  or branch-wise boundary constraints to better align growth with biology.

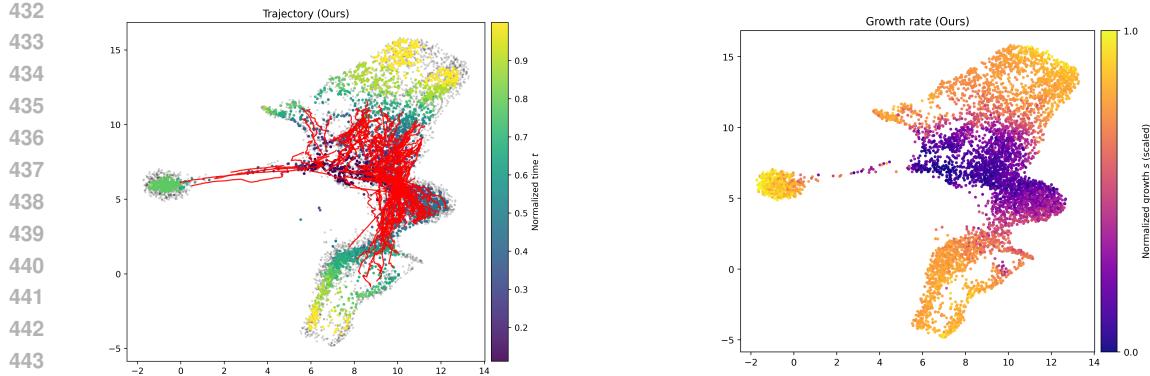


Figure 3: **Left:** Bone marrow trajectories. Colored points (ours) show inferred cell trajectories overlaid on the reference manifold (gray)) **Right:** Learned growth field — orange means high growth, purple means low growth.

## 6 DISCUSSION AND LIMITATIONS

**On the velocity prior quality and assumptions** The velocity prior  $v_{\text{prior}}$  indeed plays a constructive—but double-edged role in our method. A good prior captures coarse dynamics - reducing the learning complexity and improving on sample efficiency. A misspecified prior can bias the learned corrective field  $s_\theta(t, x)$  and slow or destabilize training. Hence, the *quality* of  $v_{\text{prior}}$  strongly influences both optimization and generalization. In practice, mild perturbation of the prior through noise, scale or misspecification are corrected by  $s_\theta$ , whereas severe misspecification such as overly large or structurally wrong drifts can bias the learned corrective flow as tested in our ablation studies Table 1. Moreover,  $v_{\text{prior}}$  does *not* require divergence-free assumption. Our dual objective explicitly includes the cross term  $\nabla_x s_\theta \cdot v_{\text{prior}}$  in the HJB residual avoiding hidden orthogonality requirements.

**Limitations** In Fig. 3 (right) we observe high growth near terminal regions, which is biologically implausible for mature or cell-cycle-exiting states. Without additional biological constraints such as cell-cycle markers, branch-terminal boundary conditions or proliferation markers, growth–transport disentanglement may remain underdetermined in some regions. A promising direction is to add weak supervision on the learned growth model to improve identifiability.

For single-cell datasets we currently use local supervision—training on adjacent pairs with a time-continuous shared network. This choice is simple and scalable and induces a globally smooth field, but it does not jointly enforce all marginals as in recent multi-marginal methods, which may limit long-range trajectory coherence. In future works, extending VP-HJF with global consistency could improve on long-range trajectory inference.

## 7 CONCLUSION

Our method decomposes the velocity field to embed domain knowledge for better energy efficiency and uses a single network to capture both unbalanced growth and transport.

486 REFERENCES  
487

- 488 Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching  
489 for generative modeling. *arXiv preprint arXiv:2210.02747*, 2022.
- 490 Michael S Albergo and Eric Vanden-Eijnden. Building normalizing flows with stochastic inter-  
491 polants. *arXiv preprint arXiv:2209.15571*, 2022.
- 492 Xingchao Liu, Chengyue Gong, and Qiang Liu. Flow straight and fast: Learning to generate and  
493 transfer data with rectified flow. *arXiv preprint arXiv:2209.03003*, 2022.
- 494 Kirill Neklyudov, Rob Brekelmans, Daniel Severo, and Alireza Makhzani. Action matching: Learn-  
495 ing stochastic dynamics from samples. In *International conference on machine learning*, pages  
496 25858–25889. PMLR, 2023a.
- 497 Luigi Ambrosio, Nicola Gigli, and Giuseppe Savaré. *Gradient flows: in metric spaces and in the*  
498 *space of probability measures*. Springer, 2005.
- 499 Aram-Alexandre Pooladian, Heli Ben-Hamu, Carles Domingo-Enrich, Brandon Amos, Yaron Lip-  
500 man, and Ricky TQ Chen. Multisample flow matching: Straightening flows with minibatch cou-  
501 plings. *arXiv preprint arXiv:2304.14772*, 2023.
- 502 Alexander Tong, Kilian Fatras, Nikolay Malkin, Guillaume Huguet, Yanlei Zhang, Jarrid Rector-  
503 Brooks, Guy Wolf, and Yoshua Bengio. Improving and generalizing flow-based generative models  
504 with minibatch optimal transport. *arXiv preprint arXiv:2302.00482*, 2023a.
- 505 Cédric Villani et al. *Optimal transport: old and new*, volume 338. Springer, 2008.
- 506 Jean-David Benamou and Yann Brenier. A computational fluid mechanics solution to the monge-  
507 kantorovich mass transfer problem. *Numerische Mathematik*, 84(3):375–393, 2000.
- 508 Lenaic Chizat, Gabriel Peyré, Bernhard Schmitzer, and François-Xavier Vialard. Unbalanced opti-  
509 mal transport: Dynamic and kantorovich formulations. *Journal of Functional Analysis*, 274(11):  
510 3090–3123, 2018.
- 511 Zhenyi Zhang, Tiejun Li, and Peijie Zhou. Learning stochastic dynamics from snapshots through  
512 regularized unbalanced optimal transport. *arXiv preprint arXiv:2410.00844*, 2024.
- 513 Dongyi Wang, Yuanwei Jiang, Zhenyi Zhang, Xiang Gu, Peijie Zhou, and Jian Sun. Joint velocity-  
514 growth flow matching for single-cell dynamics modeling. *arXiv preprint arXiv:2505.13413*, 2025.
- 515 Takeshi Koshizuka and Issei Sato. Neural lagrangian schr\” odinger bridge: Diffusion modeling for  
516 population dynamics. *arXiv preprint arXiv:2204.04853*, 2022.
- 517 Kirill Neklyudov, Rob Brekelmans, Alexander Tong, Lazar Atanackovic, Qiang Liu, and Alireza  
518 Makhzani. A computational framework for solving wasserstein lagrangian flows. *arXiv preprint*  
519 *arXiv:2310.10649*, 2023b.
- 520 Alexander Tong, Jessie Huang, Guy Wolf, David Van Dijk, and Smita Krishnaswamy. Trajectorynet:  
521 A dynamic optimal transport network for modeling cellular dynamics. In *International conference*  
522 *on machine learning*, pages 9526–9536. PMLR, 2020.
- 523 Guillaume Huguet, Daniel Sumner Magruder, Alexander Tong, Oluwadamilola Fasina, Manik  
524 Kuchroo, Guy Wolf, and Smita Krishnaswamy. Manifold interpolating optimal-transport flows  
525 for trajectory inference. *Advances in neural information processing systems*, 35:29705–29718,  
526 2022.
- 527 Alexander Tong, Nikolay Malkin, Kilian Fatras, Lazar Atanackovic, Yanlei Zhang, Guillaume  
528 Huguet, Guy Wolf, and Yoshua Bengio. Simulation-free schr\” odinger bridges via score and  
529 flow matching. *arXiv preprint arXiv:2307.03672*, 2023b.
- 530 Yuhao Sun, Zhenyi Zhang, Zihan Wang, Tiejun Li, and Peijie Zhou. Variational regularized unbal-  
531 anced optimal transport: Single network, least action. *arXiv preprint arXiv:2505.11823*, 2025.

- 540 Kevin R Moon, David Van Dijk, Zheng Wang, Scott Gigante, Daniel B Burkhardt, William S Chen,  
541 Kristina Yim, Antonia van den Elzen, Matthew J Hirn, Ronald R Coifman, et al. Visualizing  
542 structure and transitions in high-dimensional biological data. *Nature biotechnology*, 37(12):1482–  
543 1492, 2019.
- 544 Panagiotis Theodoropoulos, Augustinos D Saravacos, Evangelos A Theodorou, and Guan-  
545 Horng Liu. Momentum multi-marginal schrödinger bridge matching. *arXiv preprint arXiv:2506.10168*, 2025.
- 546 Yunyi Shen, Renato Berlinghieri, and Tamara Broderick. Multi-marginal schrödinger bridges  
547 with iterative reference refinement. *arXiv preprint arXiv:2408.06277*, 2024.
- 548 Martin Rohbeck, Edward De Brouwer, Charlotte Bunne, Jan-Christian Huetter, Anne Biton,  
549 Kelvin Y Chen, Aviv Regev, and Romain Lopez. Modeling complex system dynamics with flow  
550 matching across time and conditions. In *The Thirteenth International Conference on Learning  
551 Representations*, 2025.
- 552 Tianrong Chen, Guan-Horng Liu, Molei Tao, and Evangelos Theodorou. Deep momentum multi-  
553 marginal schrödinger bridge. *Advances in Neural Information Processing Systems*, 36:57058–  
554 57086, 2023.
- 555 Volker Bergen, Marius Lange, Stefan Peidli, F. Alexander Wolf, and Fabian J. Theis. Generalizing  
556 rna velocity to transient cell states through dynamical modeling. *Nature Biotechnology*, 38(12):  
557 1408–1414, August 2020. ISSN 1546-1696. doi: 10.1038/s41587-020-0591-3. URL <http://dx.doi.org/10.1038/s41587-020-0591-3>.
- 558  
559  
560  
561  
562  
563  
564  
565  
566  
567  
568  
569  
570  
571  
572  
573  
574  
575  
576  
577  
578  
579  
580  
581  
582  
583  
584  
585  
586  
587  
588  
589  
590  
591  
592  
593

594    **A PRIOR-GUIDED HJB INEQUALITY DERIVATION**  
 595

596    To derive the velocity-prior guided HJB residual inequality in Eq.8.

597    
$$\begin{aligned} \mathcal{A}(\rho, v, g) &\geq \mathbb{E}_{x \sim \rho_0}[s_0(x)] - \mathbb{E}_{x \sim \rho_1}[s_1(x)] \\ 598 &\quad - \int_0^1 \int \rho_t(x) \left( \partial_t s + \frac{1}{2} \|\nabla s\|^2 + \nabla s \cdot v_{\text{prior}} + \frac{1}{2} s^2 \right) dx dt. \end{aligned} \quad (16)$$

601    We want to start by minimizing Definition 3.2. Since  $v_{\text{prior}}$  is constant and known, we only need to  
 602    minimize the "learned" action portion  $v_{\text{corr}}$ , then the problem becomes the following:  
 603

604    Consider the velocity-prior guided WFR action:

605    
$$\begin{aligned} \min_{\rho, v_{\text{ot}}, g} \quad & \mathcal{A}(\rho, v_{\text{corr}}, g) := \int_0^1 \int_{\Omega} \left( \frac{1}{2} \|v_{\text{corr}}(t, x)\|^2 + \frac{1}{2} g(t, x)^2 \right) \rho_t(x) dx dt \\ 606 & \text{s.t. } \partial_t \rho_t + \nabla \cdot (\rho_t(v_{\text{prior}} + v_{\text{corr}})) = g_t \rho_t, \quad \rho|_{t=0} = \rho_0, \quad \rho|_{t=1} = \rho_1. \end{aligned} \quad (17)$$

609    **Step 1: Lagrangian formulation**

610    First, we introduce a scalar multiplier  $s(t, x)$ , the Lagrangian becomes:

611    
$$\begin{aligned} \mathcal{L} &= \int_0^1 \int_{\Omega} \left( \frac{1}{2} \|v_{\text{corr}}\|^2 + \frac{1}{2} g^2 \right) \rho dx dt \\ 612 &\quad + \int_0^1 \int_{\Omega} s \left( \partial_t \rho + \nabla \cdot (\rho(v_{\text{prior}} + v_{\text{corr}})) - g \rho \right) dx dt. \end{aligned} \quad (18)$$

616    **Step 2: Integration by parts**

617    Integration by parts in time on  $\int s \partial_t \rho$

618    
$$\begin{aligned} \int_0^1 \int_{\Omega} s \partial_t \rho dx dt &= \left[ \int_{\Omega} s \rho dx \right]_{t=0}^{t=1} - \int_0^1 \int_{\Omega} \rho \partial_t s dx dt \\ 619 &= \mathbb{E}_{x \sim \rho_0}[s_0(x)] - \mathbb{E}_{x \sim \rho_1}[s_1(x)] - \int_0^1 \int_{\Omega} \rho \partial_t s dx dt. \end{aligned} \quad (19)$$

623    Integration by parts in space on  $\int s \nabla \cdot (\rho w)$  with  $w := v_{\text{prior}} + v_{\text{ot}}$

625    
$$\int_{\Omega} s \nabla \cdot (\rho w) dx = \int_{\partial \Omega} s \rho w \cdot n d\sigma - \int_{\Omega} \nabla s \cdot (\rho w) dx. \quad (20)$$

627    Assuming zero boundary flux or fast decay, the boundary term vanishes:

628    
$$\int_0^1 \int_{\Omega} s \nabla \cdot (\rho w) dx dt = - \int_0^1 \int_{\Omega} \rho w \cdot \nabla s dx dt. \quad (21)$$

631    Then, the Lagrangian becomes

632    
$$\begin{aligned} \mathcal{L} &= \mathbb{E}_{x \sim \rho_0}[s_0(x)] - \mathbb{E}_{x \sim \rho_1}[s_1(x)] \\ 633 &\quad + \int_0^1 \int_{\Omega} \rho \left[ \frac{1}{2} \|v_{\text{corr}}\|^2 - v_{\text{corr}} \cdot \nabla s + \frac{1}{2} g^2 - s g - \partial_t s - \nabla s \cdot v_{\text{prior}} \right] dx dt. \end{aligned} \quad (22)$$

636    **Step 3: Fenchel–Young inequality**

637    The Fenchel–Young inequality states that for any vectors  $a$  and  $p$ , we have:

638    
$$\frac{1}{2} \|a\|^2 \geq p \cdot a - \frac{1}{2} \|p\|^2 \quad (23)$$

639    We set  $a = v_{\text{corr}}$  and  $p = \nabla s$ , then we have:

640    
$$\frac{1}{2} \|v_{\text{corr}}\|^2 - v_{\text{corr}} \cdot \nabla s \geq -\frac{1}{2} \|\nabla s\|^2 \quad (24)$$

641    Similarly, we set  $a = g$  and  $p = s$ , then we have:

642    
$$\frac{1}{2} g^2 - s g \geq -\frac{1}{2} s^2, \quad (25)$$

644    with equality iff  $v_{\text{ot}} = \nabla s$  and  $g = s$ . Thus, putting pieces together, we have:

645    
$$\begin{aligned} \mathcal{A}(\rho, v_{\text{corr}}, g) &\geq \mathbb{E}_{x \sim \rho_0}[s_0(x)] - \mathbb{E}_{x \sim \rho_1}[s_1(x)] \\ 646 &\quad - \int_0^1 \int \rho_t(x) \left( \partial_t s + \frac{1}{2} \|\nabla s\|^2 + \nabla s \cdot v_{\text{prior}} + \frac{1}{2} s^2 \right) dx dt. \end{aligned} \quad (26)$$

648  
649**B ADDITIONAL EXPERIMENTS**650  
651  
652  
653  
654  
655  
656  
657

**Unbalanced case - Gaussian Translation** We create a synthetic gaussian dataset to mimic the cell-fate setting by creating two uneven mass densities with a RNA velocity prior. We set  $v_{\text{prior}} = b + \eta(x - \mu_0)$ ,  $b = \mu_1 - \mu_0$  which shift the means with a mild linear drift. By prop.3.1, the learned  $s_\theta$  models the growth term  $g(x) = s_\theta(x)$  to correct the mass imbalance, while  $\nabla s_\theta(x)$  adjusts the geometric shape deformation. Fig.2 (right) displays mass-aware hexbins together with 95% covariance ellipses. Our VP-HJF closely aligns its ellipse with the target, indicating the mean and covariance agreement. The light color intensity of the bins indicate the local mass allocation match between the predicted bin masses (blue) and target (orange).

658  
659  
660  
661  
662  
663  
664  
665  
666  
667  
668

Moreover, we also examined the predicted terminal total mass  $\sum_i w_i(1)$  and the ground truth in probability. Ours matches the target unbalanced data mass ratio, whereas AM struggled and failed to adjust mass change. This highlights the need for an explicit mass constraint loss such as  $L_{\text{mass}}$ . In principle, the unbalanced AM objective in Eq. 4 does allow for mass change, since the growth term  $g_t = s_t$  governs  $\frac{d}{dt} \log w_t$  while transport is driven by  $\nabla s_t$ . In practice, however, the model often collapses to the trivial solution  $s_t \approx 0$ , which preserves the *total* mass, or becomes unstable with exponentially exploding weights. For the trivial case, suppose  $\sum_i w_i(0) = 1$ , and define  $M(t) = \int \rho_t(x) dx$ . Then  $\ln M(1) - \ln M(0) = \int_0^1 \mathbb{E}_{\rho_t}[s_t(x)] dt$ . If  $s_t \rightarrow 0$ , this implies  $M(1) \approx 1$ , consistent with our empirical observations. While such behavior can suffice for local reweighting with zero net growth (as in Neklyudov et al. (2023a), Fig. 7), it fails when the total mass must change.

669  
670  
671**C LLM USAGE**672  
673  
674  
675  
676  
677  
678  
679  
680  
681  
682  
683  
684  
685  
686  
687  
688  
689  
690  
691  
692  
693  
694  
695  
696  
697  
698  
699  
700  
701